

Galois Sync-Up

Thomas Gilray & Kris Micinski
(Task 1, POLYMORPH V-SPELLS)
April 20, 2022



High-Level Updates

- Massively-parallel deductive inference for program analysis
 - **Slog**: structured Datalog via parallel relational algebra
- Apropos Task 1: component identification via static analysis
- Progress: data-parallel binary instruction seq. matching
- Planning next steps
 - Cloud-based data-parallel deductive analytics
 - Scaling Slog to more-fully-featured Datalog

Slog: Massively Parallel Deductive Inference

- State-of-the-art program analysis: Soufflé (Datalog)
- Compiles to tight `for` loops in C++, efficient datastructures (trie, brie)
- Almost no parallelism exploited
- Our new engine (V-SPELLS): **Slog**
- Extends Datalog with *interning*, compiles to parallel RA kernels


Slog scales well—even on one node—but some sequential overhead

Graph	System	Time (s)			
		15	30	60	120
fb-media	SLOG	62	40	21	18
	Soufflé	35	33	34	37
	Radlog	254	295	340	164
ring10000	SLOG	363	218	177	115
	Soufflé	149	143	140	141
	Radlog	464	646	852	1292
suitesparse	SLOG	–	1,593	908	671
	Soufflé	1,417	1,349	1,306	1,282
	Radlog	–	–	–	–

Transitive Closure Computation
Machine: 128 threads, 20TB RAM, Azure

Strong Scaling on the Theta supercomputer

- Scaled core control-flow analysis for the lambda calculus on the Theta SC
- Promising strong scaling results up to 800 processes
- Communication overhead dominates >1k threads
- Currently innovating in this direction w/ Sidharth Kumar
- Algorithmic innovation via two-phase variadic Bruck



Task: Component Identification

- Slog achieves: parallel semantic-enabled deductive reasoning
- Goal (Task 1): **identify matching components in large codebase**
- Ported fragments of **ddisasm** (Datalog disassembly) to Slog
- Initial goal: identify equivalent instruction sequences in binaries
- Sequences of 2, 4, 8, ...
- ```
[(inst_seq8 Start1 _ Bin_name1 Ins_tree)
 (inst_seq8 Start2 _ Bin_name2 Ins_tree)
 (=/= Bin_name1 Bin_name2)
 -->
 (exact_same_sequence Bin_name1 Start1 Bin_name2 Start2 Ins_tree)]
```

# Instruction Sequence Identification—Results

- Good absolute scaling results vs. soufflé
- k-ary joins in Slog incur higher algorithmic complexity
  - Now working on strategies to mitigate this
- Memory usage in Slog was unnecessarily high
  - Rewrote our data structure to be ~2-3x vs. Souffle's optimized version



| exact sequence matching for 2 binaries ~300k |       |          |             |                                                                                                                                                                                                                              |
|----------------------------------------------|-------|----------|-------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                              | cores | time (s) | memory (MB) | other note                                                                                                                                                                                                                   |
| pure souffle                                 | 1     | 81.86    | 861         | very naive way<br><a href="https://github.com/harp-lab/reversi/blob/master/datalog/souffle/inst_seq.dl">https://github.com/harp-lab/reversi/blob/master/datalog/souffle/inst_seq.dl</a>                                      |
| souffle + ADT                                | 1     | 3.94     | 36          | output is already structure<br><a href="https://github.com/harp-lab/reversi/blob/master/datalog/souffle/inst_seq_adt.dl">https://github.com/harp-lab/reversi/blob/master/datalog/souffle/inst_seq_adt.dl</a>                 |
| slog + trie/shmap                            | 20    | 11.94    | 7000        | using `shmap.h`, actually output tuple counts is 89mb<br><a href="https://github.com/harp-lab/reversi/blob/master/datalog/slog/inst_seq.slog">https://github.com/harp-lab/reversi/blob/master/datalog/slog/inst_seq.slog</a> |
| slog + trie/google_btree                     | 20    | 4.18     | 750         | using `shmap_goog.h`                                                                                                                                                                                                         |
| slog + trie/souffle btree no cache           | 6     | 14.03    | 5000        | in docker using release mode building not work but works in debug mode                                                                                                                                                       |
| slog + google btree                          | 20    | 2.26     | 402.6       |                                                                                                                                                                                                                              |
| slog + souffle btree                         | 20    | 2.62     | 351         | Cache is not used                                                                                                                                                                                                            |

# Roadblocks / Questions for Galois

- Ordered servers from Dell
  - Chip shortage and staffing issues causing these to be delayed
  - Projecting they will be set up and running by summer
  - Will continue to rent machines from AWS / Azure to benchmark
- Possible questions:
  - Is Galois using Datalog for their performance on Task 1?
  - Potential direction—port cclyzer, to analyze LLVM code via Soufflé



# Deductive Inference on the Cloud

- We understand Slog to be the state of the art in parallel deductive inference
- In the future, we expect deductive inference to be performed in the cloud
- Want to measure: deductive inference on AWS via ParallelCloud
  - Dollar-for-dollar: Soufflé vs. Radlog vs. Slog
    - VLDB experiments and benchmark track